

SQL - Read from multiple tables

Example: relational model

student

id	name	zip
s1	jack	10027
s2	jill	11211
s3	qing	10003
s4	arun	10012

major

id	name
m1	physics
m2	biology
m3	chemistry

student-major

student_id	major_id
s1	m1
s1	m2
s2	m3
s4	m2

Example: SQL

Gather all the data for Jack

```
select * from student  
where name = 'Jack'
```

```
select major.name  
from major, student-major  
where student-major.student_id = "s1"  
    and major.id = student-major.major_id
```

Example: SQL

Get all the data for Jack in one query (join)

```
select s.name, s.zip, m.name  
from student s  
  join student-major sm on s.id = sm.student_id  
  join major m on sm.major_id = m.id  
where s.name = 'Jack'
```

Example: SQL

Get all data that we can (outer join)

```
select s.name, s.zip, m.name  
from student s  
left join student-major sm on s.id = sm.student_id  
join major m on sm.major_id = m.id
```

MongoDB version

```
{
  _id: s1,
  name: 'jack',
  zip: '10027',
  major: ['physics','biology']
}
{
  _id: s2,
  .....
}
```

MongoDB - Schema-less DB

Structure

- Documents are the smallest unit
 - JSON-like format
- Documents form a collection
 - Each document is a 'record' in a collection
 - Documents don't have to have the same format

MongoDB pros and cons

Pros

- Flexibility with Structure
- Object structure is clear
- No joins!
- Easy to scale
- Maps to application objects
- Easy to replicate

Cons

- No transaction support
- Data integrity?
- Replica lags may be a problem in write-heavy databases
- Not good for multi-user systems (“global” write locks)

Set up

The use of MongoDB is more “manual” than other database engines

- Installation

- On Mac – Using homebrew or Manually.

- Post install

- Look at the config file:

- Default location – `/usr/local/etc/mongod.conf`

- Identify locations of log file and data folder.

- Start the Mongo Server

- Using config file – `mongod --config <path-to-config-file>`

- Pointing to data location – `mongod --dbpath <path-to-data folder>`

Basics Actions

MongoDB is “forgiving” – Entities will be created if they don’t exist

- Open Shell

- `mongo` or `mongo <db name>`

- Database doesn’t have to pre-exist

- List databases

- `show databases`

- Create new database

- `use <db name>`

- Delete a database

- `db.dropDatabase()`

Basics Actions

- List collections (tables)
 - show collections
- Create new collection
 - First insert creates the collection –
(See Insert command)
- Define an index
 - `db.<collection name>.createIndex({<key>: 1}, { unique: true })`
 - 1 is ascending order. -1 is descending
- Add a new column
 - Not relevant – no columns in Mongo
- Delete a collection
 - `db.<collection name>.drop()`

CRUD Operations

→ CREATE – add documents (records) to a collection:

→ Syntax – `db.<collection name>.insert(<document>)`

→ `<document>` is a valid json – `{key1: val1, key2: { key3: val3 } }`

→ If the `_id` field is not provided – documents get an `_id` field assigned with a unique `ObjectId`.

insert example

```
db.restaurants.insert(  
  {  
    "address" : {  
      "street" : "2 Avenue",  
      "zipcode" : "10075",  
      "building" : "1480",  
      "coord" : [ -73.9557413, 40.7720266 ]  
    },  
    "borough" : "Manhattan",  
    "cuisine" : "Italian",  
    "grades" : [  
      {  
        "date" : ISODate("2014-10-01T00:00:00Z"),  
        "grade" : "A",  
        "score" : 11  
      },  
      {  
        "date" : ISODate("2014-01-16T00:00:00Z"),  
        "grade" : "B",  
        "score" : 17  
      }  
    ],  
    "name" : "Vella",  
    "restaurant_id" : "41704620"  
  }  
)
```

CRUD Operations

→ READ – find data in a collection

→ Syntax – `db.<collection name>.find(<condition>, <projection>)`

→ All Data – `db.<collection name>.find()`

→ Conditionally list from a collection.

→ `db.<collection name>.find(<condition>)`

→ *<condition> is a json document –*

→ `{}` – *returns all documents*

→ `{ key: value }`

→ `{ key: { <operator>: value } }`

→ Multiple Conditions:

→ AND – `{ key1: value1, key2: value2 }`

→ OR – `{ $or: [{key1: value1}, {key2: value2}] }`

→ Keys can be embedded fields, e.g. “grades.grade”.

→ Examples of <operator> are \$eq, \$gt, \$in

CRUD Operations

→ READ – find data in a collection

→ Conditionally get fields – `db.<collection name>.find({}, <projection>)`

→ *<projection> is a valid json listing the fields (_id is returned by default) –*

→ `{key1: 1, key2: 1}` – return only key1 and key2 (and _id)

→ `{_id: 0}` – don't return _id

→ note – true/false can replace 1/0

→ By Default mongo returns 20 documents. If there are more – type it to iterate through rest of the records.

→ To sort documents – `db.<collection name>.find().sort({key: 1})`

→ To display as indented json – `db.<collection name>.find().pretty()`

CRUD Operations

→ READ – find data in a collection

→ Many more “helper” find methods:

- `findAndModify()`
- `findOne()`
- `findOneAndDelete()`
- `findOneAndReplace()`
- `findOneAndUpdate`

CRUD Operations

→ **UPDATE** – update data in a collection

→ **Syntax** – `db.<collection name>.update(<condition>, <update>, <options>)`

→ **<condition>** – we know...

→ **<update>** – a json document with keys and values to update

→ `{key1: new_val1, key2: new_val2}` – replaces the whole document

→ `{$set: { key1: new_val1}}` – updates only the specified keys

→ **<options>** – a json document:

→ By Default mongo updates a single document. To update many –
`{multi: true}`

→ Upsert – to add a record if no match for the condition –
`{upsert: true}`

CRUD Operations

→ DELETE – delete data from a collection

→ Syntax –

→ `db.<collection name>.deleteOne(<condition>)`

→ `db.<collection name>.deleteMany(<condition>)`

→ `db.<collection name>.remove(<condition>, <options>)`

→ *<condition>* – we know...

→ *<options>* – justOne

→ Defaults to true – to limit deletion to one record

→ Set to false to allow multiple document deletion

CRUD Operations

→ Advanced READ

→ Joining collections

→ Traditionally done in two steps

→ Find documents in collection1

→ Use the values of the “join” field from collection1 to find documents in collection2

→ Use \$aggregate and \$lookup

```
→ db.collection2.aggregate([
  {$match: {
    key1: val1
  }},
  {$lookup: {
    from: "collection1",
    localField: "key2",
    foreignField: "key2",
    as "key2_values"}
  }
])
```

MongoDB in Python